

Contents

Prompt Engineering for AI Assisted Coding	2
Weak Prompt (Common Mistake)	2
Strong Prompt Structure for Cursor.....	2
Techniques for Cursor Prompting.....	3
Tell Cursor the Architecture First	3
Use "Act As" Roles	3
Use Multi-Step Prompts (Very Powerful).....	3
Use Diff-Based Instructions	4
Use "Explain Before Changing".....	4
The RTF Framework.....	4
Chain-of-Thought Prompting.....	5
Zero-shot vs. Few-shot Prompting.....	7
Cursor-Specific Tips	8
Be explicit about file scope	8
Limit hallucinations	8
Control verbosity	8
Prompt Patterns for Developers	9
Common Cursor Mistakes	9

Prompt Engineering for AI Assisted Coding

When using Cursor AI, prompt engineering is slightly different from normal ChatGPT prompting because:

- It works inside your codebase
- It has file awareness
- It can refactor, edit, and generate code directly
- Context window management matters more

Cursor works best when you give it:

- a. Clear objective
- b. Scope (which files?)
- c. Constraints (framework, version, style)
- d. Expected output (refactor? explanation? patch?)

Weak Prompt (Common Mistake)

Fix this code.

Why is this weak:

- What kind of fix?
- Refactor or debug?
- Keep architecture?
- Improve performance?
- Follow which standard?

Strong Prompt Structure for Cursor

Use this template:

Context:

This is a Spring Boot 3 REST API project using Java 21.

Goal:

Refactor this controller to follow best practices.

Constraints:

- *Keep current endpoint paths unchanged*
- *Use constructor injection*
- *Apply proper exception handling*
- *Follow clean code principles*

Output:

Modify the file directly and explain major improvements.

Techniques for Cursor Prompting

Tell Cursor the Architecture First

Cursor performs better when you declare architecture.

Example:

This project follows:

- Layered architecture
- Controller → Service → Repository pattern
- JPA with PostgreSQL
- DTO-based request/response

Update the service layer to improve transaction handling.

Without this, it may invent structure.

Use “Act As” Roles

Example:

Act as a senior backend engineer reviewing this code.

Identify:

- Security risks
- Performance bottlenecks
- Violations of SOLID principles

This improves analysis depth dramatically.

Use Multi-Step Prompts (Very Powerful)

Instead of:

Add authentication.

Do this:

Step 1: Analyze current security setup.

Step 2: Suggest best authentication approach for this project.

Step 3: Implement JWT-based authentication.

Step 4: Add necessary configuration classes.

You get controlled evolution instead of chaotic generation.

Use Diff-Based Instructions

Cursor is excellent at controlled modifications.

Instead of:

Improve this function.

Try:

Refactor this function but:

- Do not change its signature
- Do not modify database schema
- Keep backward compatibility
- Reduce cognitive complexity

This prevents over-engineering.

Use "Explain Before Changing"

Very powerful for learning:

Before modifying the code:

1. Explain what is wrong
2. Propose 2 solutions
3. Then implement the best one

This turns Cursor into a mentor instead of just a generator.

The RTF Framework

To get professional results, you must treat the AI like a specialized colleague. Using the **RTF (Role, Task, Format)** framework prevents generic or "lazy" code.

- **Role:** Define the expertise level (e.g., "Senior Python Developer," "Security Auditor").
- **Task:** Define the specific action (e.g., "Refactor this loop," "Write unit tests").
- **Format:** Define how you want the answer (e.g., "A Git Diff," "A Markdown table," "Commented code").

Example Comparison:

- **Weak Prompt:** "Clean up this code."
- **RTF Prompt:** "Act as a **Senior React Developer**. **Refactor** the following component to use the useReducer hook instead of multiple useState calls. Output the result as a **Side-by-Side Diff**."

Scenario 1: A messy, "junior-level" JavaScript function that needs a professional refactor.

calculator.js:

```
function calc(a, b, type) {  
  if (type == "add") {  
    return a + b;  
  } else if (type == "sub") {  
    return a - b;  
  } else {  
    return "error";  
  }  
}
```

The Prompt to Demo:

"Act as a **Senior TypeScript Developer**. **Refactor** this function to use an enum for the operation types and ensure the function is strictly typed. **Output** the result as a modern TypeScript file."

Scenario 2: The "Logic & Security" Lab

Practice **RTF Prompting** to fix security flaws and logic gaps in a banking context.

Open this code and use Ctrl + K. Act as a **Senior Security Auditor** and ask the AI to fix the vulnerabilities.

```
// File: bankTransfer.js  
function transferFunds(amount, accountId) {  
  // BUG: No validation if amount is negative  
  // BUG: No check if accountId exists  
  console.log("Transferring " + amount + " to " + accountId);  
  
  let balance = 1000;  
  balance = balance - amount;  
  
  return "Success! New balance: " + balance;  
}
```

Recommended Prompt:

"Act as a **Senior Security Auditor**. **Refactor** this transfer function to: 1. Prevent negative transfers, 2. Ensure the accountId is a 10-digit string, and 3. Check for sufficient balance. **Output** the result as clean, production-ready code."

Chain-of-Thought Prompting

AI models often fail when they try to write 100 lines of code at once. **Chain-of-Thought** prompting forces the AI to "think out loud" and create a logical plan before it touches the keyboard. This significantly reduces "hallucinated" libraries or broken logic.

Step-by-Step Demo: The "Plan-First" Approach

1. Open Chat (Ctrl + L).
2. Input a complex request but end it with a planning instruction:

*"I need a Node.js script that connects to a SQL database and exports transactions to a CSV. **Before writing any code, list the steps and logic you will follow. Wait for my approval.**"*

3. Review the AI's plan. If a step is wrong (e.g., it suggests the wrong library), correct it.
4. Type "**Proceed**" to have it generate the code based on the approved plan.

Scenario 1: Building a complex logic flow for a banking application (matching your interest in the financial sector).

The Prompt to Demo (In Ctrl + L Chat):

*"I need to build a 'Surprise Assumption of Duties' tracker for a bank. It needs to check if a manager has been on mandatory leave for 5 days, and if so, trigger a notification for a 'Surprise Assumption' by a relief officer. **Before writing any code, explain the logic flow and the database schema I would need. Do not write the code yet.**"*

Scenario 2: The "API & Async" Lab

Practice **Chain-of-Thought (CoT)** to handle asynchronous data fetching properly.

This code crashes because it doesn't wait for the data. Use Ctrl + L to ask the AI to **plan** the fix before writing it.

```
// File: userData.js
function getUserProfile(userId) {
  const data = fetch(`https://api.bank.com/users/${userId}`);
  // BUG: 'fetch' is async, but this code treats it as sync
  console.log(data.name);
  return data;
}
```

Recommended Prompt:

*"This code is failing because of asynchronous timing issues. **Before fixing it, explain the steps** needed to implement a proper async/await pattern with a try/catch block for error handling. Once I approve the plan, write the code."*

Zero-shot vs. Few-shot Prompting

Sometimes you want the AI to follow your specific coding "vibe"—like how you name variables or how you handle errors.

- **Zero-shot:** You ask for code with no examples. The AI uses its default style.
- **Few-shot:** You provide 1 or 2 examples of your preferred style first, then ask for the task.

Few-shot Demonstration:

1. Open **Inline Edit (Ctrl + K)**.
2. Provide context by pointing to your existing style:

"Here is how I handle errors in this project: try { ... } catch (e) { logError(e); }. Now, using that same style, write a function that fetches user data from @api.js."

3. Observe how the AI mirrors your exact error-handling pattern instead of using a generic console.log.

Scenario 1: Teaching the AI to follow a specific, custom error-handling pattern used in your classroom or company.

Copy this "Pattern Style" into a file named `style_guide.js`:

```
// OUR PROTECTED PATTERN:  
// We always wrap in try-catch and use the custom 'BankLogger' class.  
try {  
    // logic here  
} catch (err) {  
    BankLogger.error("Process Failed", { timestamp: Date.now(), details: err });  
}
```

The Prompt to Demo (In Ctrl + K while highlighting the code):

*"Using the **same pattern** shown here, write a new function called `processLoanApplication` that takes a user object and sends it to an API endpoint /v1/apply."*

Scenario 2: The "Style & Consistency" Lab

Practice **Few-shot Prompting** to match a specific UI component style.

Look at the "Pattern" component below. Use Ctrl + K to generate a NEW component (a `NotificationCard`) that looks exactly like it.

```
<div class="card shadow-sm border-blue">  
  <h3 class="text-bold">Transaction Alert</h3>  
  <p class="text-muted">Please review your recent activity.</p>  
  <button class="btn-primary">View Details</button>  
</div>
```

Recommended Prompt:

"Using the **same HTML structure and CSS classes** (`shadow-sm`, `text-bold`, `btn-primary`) as the Transaction Alert card above, create a new NotificationCard. It should have a title for 'System Update' and a button that says 'Update Now!'"

Cursor-Specific Tips

Be explicit about file scope

Instead of vague requests:

Improve the project.

Say:

Focus only on:

- *UserController.java*
- *UserService.java*

Ignore other modules.

Limit hallucinations

If you want strict edits:

- Do not create new classes unless absolutely necessary.
- Do not introduce new dependencies.

Control verbosity

Keep explanations short.

Only explain architectural changes.

or

Explain in detail because I am learning.

Prompt Patterns for Developers

Project Context:

[Framework, version, database, architecture]

Current Problem:

[Describe clearly]

Objective:

[What you want improved]

Constraints:

- No breaking changes
- Maintain API contract
- Use best practices for [framework]

Quality Requirements:

- Clean code
- SOLID principles
- Production-ready

Output:

- Show diff-style changes
- Explain reasoning

This works incredibly well in Cursor.

Common Cursor Mistakes

- Giving vague commands
- Asking for too many changes at once
- Not specifying constraints
- Not defining architecture
- Letting it rewrite too much